

```
In [ ]: import os
```

```
In [ ]: # MDP Environment
import random

class MDP(object):

    def __init__(self, Questions, Prob_Rewards):
        self.Questions = Questions
        self.Prob_Rewards = Prob_Rewards

    def InitState(self):
        return 1
    def EndState(self, state):
        return state == self.Questions
    def Actions(self, state):
        if self.EndState(state):
            return []
        else:
            return ['Stay', 'Quit']
    def SuccProbReward(self, state, action):
        if action == 'Stay':
            if state > self.Questions:
                return [(state, 1., 0.)]
            else:
                Probability = self.Prob_Rewards[state][0]
                Reward = self.Prob_Rewards[state][1]
                return [(state+1, Probability, Reward), (self.Questions+1, 1.-Probability, 0.)]
        elif action == 'Quit':
            if state > self.Questions:
                return [(state, 1., 0.)]
            else:
                Probability = 1.0
                Reward = self.Prob_Rewards[state][1]
                return [(self.Questions+1, Probability, Reward)]
    def discount(self):
        return 0.9
    def states(self):
        return range(1, self.Questions+2)

# Value Iteration for the above Defined MDP environment

def valueIteration(mdp, epsilon=0.001):
    # initialize values of all states to zero
    V = {s: 0.0 for s in mdp.states()}
    # repeat until convergence
    while True:
        # set change to zero
        delta = 0
        # for each state, update its value using Bellman update
        for s in mdp.states():
            if mdp.EndState(s):
                continue
            # compute the maximum value over all possible Actions
            MaxValue = -float('inf')
            for a in mdp.Actions(s):
                val = 0
                for NextState, Probability, Reward in mdp.SuccProbReward(s, a):
                    val += Probability * (Reward + mdp.discount() * V[NextState])
                MaxValue = max(MaxValue, val)
            # update value of state
```

```

        delta = max(delta, abs(MaxValue - V[s]))
        V[s] = MaxValue
    # check for convergence
    if delta < epsilon:
        break
    # compute the optimal policy using the computed values
    pi = {}
    for s in mdp.states():
        if mdp.EndState(s):
            pi[s] = None
        else:
            MaxValue = -float('inf')
            best_action = None
            for a in mdp.Actions(s):
                val = 0
                for NextState, Probability, Reward in mdp.SuccProbReward(s, a):
                    val += Probability * (Reward + mdp.discount() * V[NextState])
                if val > MaxValue:
                    MaxValue = val
                    best_action = a
            pi[s] = best_action
    # return the computed values and policy
    return V, pi

Questions = 10
Prob_Rewards = {1: (0.99, 100),
                 2: (0.9, 500),
                 3: (0.8, 1000),
                 4: (0.7, 5000),
                 5: (0.6, 10000),
                 6: (0.5, 50000),
                 7: (0.4, 100000),
                 8: (0.3, 500000),
                 9: (0.2, 1000000),
                 10: (0.1, 5000000)}

kbc_MDP = MDP(Questions, Prob_Rewards)

V, P = valueIteration(kbc_MDP)
# kbc_MDP.MonteCarlo(2)
print(V)
print(P)

```

```

{1: 26780.771381760005, 2: 29945.871360000001, 3: 36414.656, 4: 49464.799999999996,
5: 72960.0, 6: 124000.0, 7: 220000.0, 8: 500000.0, 9: 1000000.0, 10: 0.0, 11: 0.0}
{1: 'Stay', 2: 'Stay', 3: 'Stay', 4: 'Stay', 5: 'Stay', 6: 'Stay', 7: 'Stay', 8:
'Quit', 9: 'Quit', 10: None, 11: 'Stay'}

```

In []: *# Policy Iteration for the above Defined MDP environment*

```

def policyIteration(mdp, epsilon=0.001):
    # initialize values of all states to zero
    V = {s: 0.0 for s in mdp.states()}
    # initialize policy arbitrarily
    pi = {s: mdp.Actions(s)[0] for s in mdp.states() if not mdp.EndState(s)}
    while True:
        # policy evaluation step
        while True:
            delta = 0
            for s in mdp.states():
                if mdp.EndState(s):
                    continue
                # compute the value of the current policy for this state
                val = 0
                for NextState, Probability, Reward in mdp.SuccProbReward(s, pi[s])

```

```

        val += Probability * (Reward + mdp.discount() * V[NextState])
        # update value of state
        delta = max(delta, abs(val - V[s]))
        V[s] = val
        if delta < epsilon:
            break
    # policy improvement step
    policy_stable = True
    for s in mdp.states():
        if mdp.EndState(s):
            continue
        old_action = pi[s]
        # find the best action using the updated values
        MaxValue = -float('inf')
        best_action = None
        for a in mdp.Actions(s):
            val = 0
            for NextState, Probability, Reward in mdp.SuccProbReward(s, a):
                val += Probability * (Reward + mdp.discount() * V[NextState])
            if val > MaxValue:
                MaxValue = val
                best_action = a
        # update policy
        pi[s] = best_action
        if old_action != best_action:
            policy_stable = False
    # check for convergence
    if policy_stable:
        break
    # return the computed values and policy
    return V, pi

Questions = 10
Prob_Rewards = {1: (0.99, 100),
                 2: (0.9, 500),
                 3: (0.8, 1000),
                 4: (0.7, 5000),
                 5: (0.6, 10000),
                 6: (0.5, 50000),
                 7: (0.4, 100000),
                 8: (0.3, 500000),
                 9: (0.2, 1000000),
                 10: (0.1, 5000000)}

kbc_MDP = MDP(Questions, Prob_Rewards)

V, P = policyIteration(kbc_MDP)

print(P)
print(V)

```

```

{1: 'Stay', 2: 'Stay', 3: 'Stay', 4: 'Stay', 5: 'Stay', 6: 'Stay', 7: 'Stay', 8:
'Quit', 9: 'Quit', 11: 'Stay'}
{1: 26780.771381760005, 2: 29945.87136000001, 3: 36414.656, 4: 49464.799999999996,
5: 72960.0, 6: 124000.0, 7: 220000.0, 8: 500000.0, 9: 1000000.0, 10: 0.0, 11: 0.0}

```

```

In [ ]: # Sate
def generateSARS(kbc_MDP, pi, num_sequences=1, max_steps=100):
    sequences = []
    for i in range(num_sequences):
        state = kbc_MDP.InitState()
        sequence = []
        for t in range(max_steps):
            if kbc_MDP.EndState(state):
                break

```



```

In [ ]: # Monto Carlo Environment for the above problem
import random

class MontoCarlo(object):

    def __init__(self, Questions, Prob_Rewards):
        self.Questions = Questions
        self.Prob_Rewards = Prob_Rewards

    def play(self, policy):
        # initialize variables
        state = self.InitState()
        done = False
        rewards = []
        # play until the game is over
        while not done:
            # choose an action according to the given policy
            action = policy(state)
            # take the action and observe the next state and reward
            next_state, reward, done = self.SuccProbReward(state, action)[0]
            # store the reward
            rewards.append(reward)
            # update the current state
            state = next_state
        # compute the cumulative rewards for each time step
        cumulative_rewards = [sum(rewards[i:]) for i in range(len(rewards))]
        # create a list of (state, action, cumulative reward) tuples
        episodes = [(self.InitState(), None, 0)]
        for i in range(len(rewards)):
            episodes.append((i+1, policy(i+1), cumulative_rewards[i]))
        return episodes

    def InitState(self):
        return 1

    def EndState(self, state):
        return state == self.Questions + 1

    def Actions(self, state):
        if self.EndState(state):
            return []
        else:
            return ['Stay', 'Quit']

    def SuccProbReward(self, state, action):
        if action == 'Stay':
            if state > self.Questions:
                return [(state, 0, True)]
            else:
                prob, reward = self.Prob_Rewards[state]
                if random.random() <= prob:
                    return [(state+1, reward, False)]
                else:
                    return [(self.Questions+1, 0, True)]
        elif action == 'Quit':
            if state > self.Questions:
                return [(state, 0, True)]
            else:
                reward = self.Prob_Rewards[state][1]
                return [(self.Questions+1, reward, True)]

    def discount(self):
        return 0.9

```

```

def states(self):
    return range(1, self.Questions+2)

# define a random policy
def random_policy(state):
    actions = ['Stay', 'Quit']
    return random.choice(actions)

# create a MontoCarlo environment
env = MontoCarlo(10, [(0.99, 100),(0.9, 500),(0.8, 1000),(0.7, 5000), (0.6, 10000)]

# run a random policy for one episode
episode = env.play(random_policy)

# print the episode
print(episode)

```

```

[(1, None, 0), (1, 'Stay', 6500), (2, 'Quit', 6000), (3, 'Quit', 5000), (4, 'Quit', 0)]

```

```

In [ ]: import numpy as np

# Define the KBC game
Questions = 10
Prob_Rewards = [(0.99, 100),(0.9, 500),(0.8, 1000),(0.7, 5000), (0.6, 10000), (0.5

# # Create the transition matrix
num_states = Questions + 1
num_actions = 2

R = np.zeros(num_states).astype(int)
for s in range(num_states):
    if s == num_states - 1:
        R[s] = 0 # end of game
    else:
        R[s] = Prob_Rewards[s][1]

# # Format the output of the transition matrix

P1 = np.array([[0.99, 0.01, 0, 0, 0, 0, 0, 0, 0, 0, 0],
               [0, 0.9, 0.1, 0, 0, 0, 0, 0, 0, 0, 0],
               [0, 0, 0.8, 0.2, 0, 0, 0, 0, 0, 0, 0],
               [0, 0, 0, 0.7, 0.3, 0, 0, 0, 0, 0, 0],
               [0, 0, 0, 0, 0.6, 0.4, 0, 0, 0, 0, 0],
               [0, 0, 0, 0, 0, 0.5, 0.5, 0, 0, 0, 0],
               [0, 0, 0, 0, 0, 0, 0.4, 0.6, 0, 0, 0],
               [0, 0, 0, 0, 0, 0, 0, 0.3, 0.7, 0, 0],
               [0, 0, 0, 0, 0, 0, 0, 0, 0.2, 0.8, 0],
               [0, 0, 0, 0, 0, 0, 0, 0, 0, 0.1, 0.9],
               [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1]])

print("Transition matrix : \n")
for i in range(num_states):
    for j in range(num_states):
        print(P1[i][j], end = " ")
    print("\n")
print("Rewards Vector : \n")
print(R)

# Set up the discount factor
gamma = 0.9

```

```

# Solve the system of linear equations
A = np.eye(num_states) - gamma * P1
b = R
V = np.linalg.inv(A).dot(b) # value function

# Print the value function vector
print("\n Value function : \n")
print(V)

```

Transition matrix :

```

0.99  0.01  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
0.0   0.9   0.1  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
0.0   0.0   0.8  0.2  0.0  0.0  0.0  0.0  0.0  0.0  0.0
0.0   0.0   0.0  0.7  0.3  0.0  0.0  0.0  0.0  0.0  0.0
0.0   0.0   0.0  0.0  0.6  0.4  0.0  0.0  0.0  0.0  0.0
0.0   0.0   0.0  0.0  0.0  0.5  0.5  0.0  0.0  0.0  0.0
0.0   0.0   0.0  0.0  0.0  0.0  0.4  0.6  0.0  0.0  0.0
0.0   0.0   0.0  0.0  0.0  0.0  0.0  0.3  0.7  0.0  0.0
0.0   0.0   0.0  0.0  0.0  0.0  0.0  0.0  0.2  0.8  0.0
0.0   0.0   0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.1  0.9
0.0   0.0   0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  1.0

```

Rewards Vector :

```

[   100    500   1000   5000  10000  50000 100000 500000 1000000
 5000000      0]

```

Value function :

```

[ 63648.10943705  759738.21429313 1598336.23017438 2480745.24693792
 3381021.26432234 4292416.05996743 5135175.18440464 5900948.36670179
 6043956.04395604 5494505.49450549      0.          ]

```

In []: